

Decoder-only conditional latent models with temporal mixture priors enable zero-shot cross-lab mouse sleep staging and substage discovery

Oliver Rosbæk Elmgreen^{1,2*}, Morten Mørup¹, Birgitte Rahbek Kornum^{2,3}

1 Section for Cognitive Systems, Department of Applied Mathematics and Computer Science, Technical University of Denmark, Kongens Lyngby, Denmark

2 Department of Neuroscience, University of Copenhagen, Copenhagen, Denmark

3 Department of Drug Design and Pharmacology, University of Copenhagen, Copenhagen, Denmark

* orel@dtu.dk

Abstract

Manual scoring of mouse sleep from electroencephalography (EEG) and electromyography (EMG) lacks consensus across laboratories, and supervised deep learning models trained at one site often fail on data from others until retrained on multi-laboratory labels. We present an unsupervised computational framework for population-level sleep staging on the Mouse Sleep Staging Validation dataset (MSSV; OpenNeuro ds006366), comparing a controlled model ladder on identical features and evaluation protocols: feature-domain hidden Markov model (HMM), HMM–Gaussian mixture variational autoencoder (HMMGMVAE) without subject conditioning, decoder-only conditional GMVAE (cGMVAE), and decoder-only conditional HMM–GMVAE (cHMM–GMVAE) with a warm sticky HMM–Gaussian mixture model (GMM) latent prior. Models are trained on leave-mice-out cross-validation folds pooling laboratories 2, 3, and 5. Held-out mice are staged by zero-shot inference: spectral epochs are encoded without subject identifiers and states are assigned from the learned mixture prior only, without per-mouse manual calibration. On joint cross-lab holdout, cHMM–GMVAE achieves best-of-three prior normalized mutual information (NMI) up to 0.65, compared with 0.28–0.41 for matched cGMVAE and 0.51 for a feature HMM baseline. A substage resolution sweep and physiology panels refine wake, NREM, and REM structure and reveal transition substates at macro boundaries. These results show that decoder-only conditioning with temporal mixture priors improves unsupervised generalization across mice and acquisition sites while supporting biologically interpretable substage discovery.

Author summary

We study how *Mus musculus* sleep can be classified automatically from brain and muscle signals when recordings come from different laboratories with different equipment. Expert human scorers disagree across labs, and supervised computer models often fail on new sites. We built a family of unsupervised models that learn shared sleep patterns from many mice, then test them on held-out animals they have never seen—without asking a human to calibrate each new recording. Our best model combines a neural network that reads EEG and EMG with a temporal prior that enforces realistic sleep bout structure. It outperforms simpler models on cross-lab holdout and discovers finer sleep substates whose brain-wave profiles match known

physiology. We validate substages with power spectra, muscle activity, and bout lengths, with guidance from sleep neuroscientists. This work offers a label-free route to comparing sleep across studies and may help researchers find transition states between wake, NREM, and REM that are hard to score by eye alone.

Introduction

Rodent sleep is scored from EEG and EMG into wake, non-rapid eye movement (NREM), and rapid eye movement (REM), yet montage and laboratory protocols differ across studies [1]. Rose et al. show that manual scoring lacks cross-lab consensus and that published deep models generalize poorly across MSSV sites until retrained on multi-laboratory data [2], motivating label-agnostic staging and rigorous holdout evaluation. Supervised models such as SPINDLE and REST achieve high cross-lab accuracy but require expert labels [3, 4]. Unsupervised tools include classical FASTER clustering [5], FASTER2 Gaussian HMMs on hand-crafted features [6], SegWay’s per-mouse k-means plus HMM smoothing [7], mcRBM substage discovery [8], recent WUCSS light/deep NREM states [9], and transition-focused supervised classifiers [10]; none combine deep latent learning, sticky HMM–GMM priors, and population zero-shot holdout on MSSV.

Our MSc thesis showed that subject-conditioned GMVAEs improve population staging ($\text{NMI} \approx 0.70$) over feature HMMs (≈ 0.51) [11], with interpretable substages but encoder conditioning that leaks subject identity at evaluation. We make three contributions: (i) decoder-only conditional variational autoencoder (cVAE) training with zero-shot holdout inference (no subject embeddings at staging); (ii) a fair model ladder ending in cHMM–GMVAE with temporal mixture priors [12–14]; and (iii) substage biology with expert validation. Synthetic validation, encoder+decoder ablations, and full $K=13$ taxonomy are in S4 Fig and S1 Appendix.

Results

Cross-lab holdout ladder

We evaluated a locked model ladder under joint population training with leave-mice-out four-fold cross-validation across MSSV laboratories 2, 3, and 5 (Fig 1). Fig 1A summarizes the cohort split; Fig 1B contrasts training (decoder with subject embedding) with zero-shot holdout staging (encoder and mixture prior only, no subject identifier); Fig 1C lists ladder rungs. Fig 2 and Table 1 report best-of-three prior NMI on held-out mice.

Fig 1. Study overview and zero-shot holdout protocol. A: MSSV multi-laboratory cohort with four-fold cross-validation; mice from three folds train a population model evaluated on the held-out fourth fold. B: Training uses decoder-only subject conditioning for reconstruction; holdout staging encodes EEG/EMG without subject identifiers and assigns states from the HMM–GMM prior only (contrast with per-mouse calibration in AccuSleep [15]). C: Model ladder from feature HMM through HMMGMVAE, cGMVAE, and cHMM–GMVAE.

On completed joint folds, cHMM–GMVAE exceeded cGMVAE by 0.11–0.35 NMI (joint mean 0.55 vs. 0.33 over three to four folds). Within-lab holdout showed laboratory heterogeneity: cHMM–GMVAE led on laboratory 3 (up to 0.746 NMI) while cGMVAE remained competitive on laboratory 2 (S2 Fig, S2 Table). HMMGMVAE locked runs were pending at manuscript preparation; per-fold breakdown is in S1 Fig.

Fig 2. Joint cross-lab holdout prior NMI. Bars show mean best-of-three prior NMI across completed folds for each ladder rung; dashed line marks thesis HMM feature baseline (0.51). HMMGMVAE bar pending locked holdout completion.

Substages and transitions

We swept mixture component count $K \in \{3, 5, 7, 9, 11, 13, 15\}$ on locked cHMM–GMVAE (joint fold 4) and selected K by validation score and cross-seed stability (S3 Fig). Fig 3 shows a compact physiology grid for discovered substages (EEG/EMG power spectral density, θ/δ ratio, EMG power, bout length, macro-label composition, subject mixing), following the mcRBM paradigm [8]. Rare or confounded states will be excluded after expert review (see S1 Appendix). Fig 4 shows the substage transition matrix and a two-hour hypnogram excerpt; transition-rich states concentrate at wake–NREM and NREM–REM boundaries [10, 16].

Table 1. Joint holdout prior NMI by fold (best-of-three seeds, locked recipes).

Fold	HMM [†]	HMMGMVAE	cGMVAE	cHMM-GMVAE
1	0.51	pending	0.361	0.612
2	0.51	pending	0.275	0.381
3	0.51	pending	0.407	pending
4	0.51	pending	0.296	0.646

[†]Thesis population leave-one-subject-out reference (feature HMM), not re-estimated per fold.

Fig 3. Substage physiology panels. Rows: discovered substates after expert curation; columns: EEG PSD, EMG PSD, θ/δ , EMG power, bout length, macro-label mix, subject mixing. Generated from joint holdout cHMM-GMVAE fold 4 until K sweep completes.

Fig 4. Substage dynamics. A: Row-normalized transition matrix from the learned sticky HMM-GMM prior. B: Two-hour hypnogram excerpt with substage color strip; vertical markers indicate macro-boundary transitions.

Table 2. Illustrative substage summary (to be finalized after expert interview).

Substage	Dominant macro	Mean bout (s)	θ/δ	Expert label
S1	Wake	TBD	TBD	active wake
S2	NREM	TBD	TBD	light NREM
S3	NREM	TBD	TBD	deep NREM
S4	transition	TBD	TBD	wake-NREM
S5	transition	TBD	TBD	NREM-REM

TBD: populate from postprocess after Birgitte interview (docs/paper/birgitte.interview_guide.md).

Discussion

Rose et al. establish that cross-lab supervised staging inherits rater disagreement [2]; our unsupervised ladder complements this by discovering structure without training on one laboratory’s label biases. Zero-shot holdout inference (encoder plus temporal prior, no subject identifier) contrasts with transductive calibration in AccuSleep and per-animal fitting in SegWay [7, 15]. Temporal mixture priors helped when bout structure and laboratory-specific spectra aligned with sequence modelling (laboratories 3 and 5); static GMM decoders sufficed on some montages (laboratory 2). Decoder-only conditioning is the main methodological advance over the thesis encoder+decoder setup for population holdout [11]. Absolute holdout NMI remains below curated thesis leave-one-subject-out performance; relative gains over the ladder and substage biology are the primary claims. Limitations include seed fragility, per-laboratory spectral front-end locks (not per-mouse manual scoring), and pending HMMGMVAE locked runs.

Materials and methods

Animals and dataset

We used the MSSV dataset (OpenNeuro ds006366) [1]: 92 mice, laboratories 2, 3, and 5, 4 s epochs, expert wake/NREM/REM labels for evaluation only. All experiments used previously published data; no new animal procedures were performed in this study.

Preprocessing and features

Per-laboratory spectral front-end locks (`cvae_overrides`) were fixed from incohort tuning [11]. Wide multilayer perceptron encoder-decoder with latent dimension 6; convolutional encoder on spectral inputs.

Model ladder

Locked fair-comparison recipes (`scripts/cv4fold/locked_recipes.py`): feature HMM baseline (thesis); HMMGMVAE ($emb=0$, sequence length $T=64$, `hmm_gmm` prior); cGMVAE (decoder-only, $emb=8$, GMM prior); cHMM-GMVAE (+ warm sticky HMM-GMM, $\kappa=0.92$) [12, 13, 17].

Training

Population training on cross-validation training folds (joint: all laboratories; within-lab: single laboratory). Subject embeddings ($emb=8$) were concatenated to the decoder only (`decoder_only_conditioning`), not the encoder, during reconstruction and prior learning.

Inference (holdout)

For held-out validation mice, sleep-state labels were obtained by encoding spectral epochs without subject identifiers, then assigning states via the learned GMM or sticky HMM-GMM prior on latent means (prior prediction). Subject embeddings were not used at holdout evaluation; the decoder was not invoked for staging.

Evaluation

Leave-mice-out four-fold holdout; prior-prediction NMI versus expert macro labels; best of three random seeds per configuration.

Substage sweep and physiology

Mixture component sweep $K \in \{3, 5, 7, 9, 11, 13, 15\}$; physiology panels via `scripts/substage_analysis/`; transition matrix from the learned prior.

Data availability

MSSV is available on OpenNeuro (ds006366). Processed features and fold manifests are described in the SPA repository.

Code availability

Training, evaluation, and figure scripts are in the SPA repository (`/work3/s204070/SPA`). A Zenodo archive with tagged release will be deposited before publication (DOI pending).

Acknowledgments

We thank collaborators in the Kornum laboratory for biological interpretation guidance.

Supporting information

S1 Fig. Per-fold holdout ladder. Grouped bar chart of prior NMI for all four folds and four model rungs (from `plot_ladder_figure.py --by-fold`).

S2 Fig. Within-laboratory holdout heatmap. Laboratory $\times$ model prior NMI for within-lab scope (do not compare absolute values to joint scope).	84 85
S3 Fig. Substage resolution sweep. Validation score versus mixture component count K on locked cHMM-GMVAE (joint fold 4).	86 87
S4 Fig. Thesis cGMVAE $K=13$ taxonomy. Full PCA, transition graph, and Fig-27-style panel from the MSc thesis.	88 89
S1 Table. Extended literature comparison. Supervised ceilings, transductive unsupervised, and static unsupervised methods versus this work.	90 91
S2 Table. Within-lab holdout prior NMI. Full table from <code>holdout_ladder.csv</code> .	92
S1 File. Locked YAML recipes. Joint and within-lab holdout configuration files.	93
S1 Appendix. Thesis supplement material. Synthetic HMM validation, raw time-domain HMM failure, decoder-only ablations, preprocessing locks, and Birgitte interview substage names.	94 95

References

1. Rose L, Zahid AN, Ciudad JG, et al. Mouse sleep staging validation dataset (MSSV). OpenNeuro. 2025. Dataset ds006366; cite full author list from dataset paper.
2. Rose L, Zahid AN, García Ciudad J, Egebjerg C, Piilgaard L, Sørensen FL, et al. Lack of Consensus for Manual Mouse Sleep Scoring Limits Implementation of Automatic Deep Learning Models. bioRxiv. 2026. Available from: <https://www.biorxiv.org/content/10.64898/2026.03.27.714381v1>. doi:10.64898/2026.03.27.714381.
3. Miladinović S, Muheim C, Bauer S, et al. SPINDLE: End-to-end learning from EEG/EMG to extrapolate animal sleep scoring across experimental settings, labs and species. PLOS Computational Biology. 2019;15(4):e1006968. doi:10.1371/journal.pcbi.1006968.
4. Wang J, et al. REST: A deep learning tool for automated mouse sleep stage classification. Sleep Advances. 2026;7(2):zpag044. Supervised Transformer ceiling; replace author list from publisher. doi:10.1093/sleepadvances/zpag044.
5. Sunagawa G, Seo M, Sato C, et al. Fasting promotes sleep through orexin neurons. Genes to Cells. 2013;18(11):1107-17. FASTER staging tool lineage.
6. Yamada RG, Matsuzawa K, Ode KL, Ueda HR. An automated sleep staging tool based on simple statistical features of mice electroencephalography (EEG) and electromyography (EMG) data. European Journal of Neuroscience. 2024;60(7):5467-86. doi:10.1111/ejn.16465.
7. Yaghoubi F, Sunderam S. SegWay: A novel sleep stage scoring system for mice and rats. MethodsX. 2016;3:369-75. doi:10.1016/j.mex.2016.03.009.
8. Katsageorgiou VM, Sona D, Zanotto M, Lassi G, Garcia-Garcia C, Tucci V, et al. A novel unsupervised analysis of electrophysiological signals reveals new sleep substages in mice. PLOS Biology. 2018;16(5):e2003663. doi:10.1371/journal.pbio.2003663.
9. Cusinato R, Gross CT, et al. Workflow for unsupervised classification of sleep states (WUCSS) in mice. Journal of Neuroscience Methods. 2024. Unsupervised light/deep NREM; replace full author list from publisher.

10. authors S. Somnotate: A probabilistic sleep stage classifier for studying vigilance state transitions. PLOS Computational Biology. 2024. Replace author list from publisher page before submission. doi:10.1371/journal.pcbi.1011793.
11. Elmgreen OR, Bigom AH. Characterizing sleep patterns by use of dynamic unsupervised machine learning approaches. MSc thesis, Technical University of Denmark. 2026. Population cGMVAE NMI $\approx$ 0.70; HMM feature baseline $\approx$ 0.51.
12. Dilokthanakul N, Mediano PAM, Garnelo M, et al. Deep unsupervised clustering with Gaussian mixture variational autoencoders. In: arXiv:1611.02648; 2016. .
13. Sohn K, Lee H, Yan X. Learning structured output representation using deep conditional generative models. In: Advances in Neural Information Processing Systems. vol. 28; 2015. .
14. Johnson MJ, Duvenaud DK, Wiltchko A, Datta SR, Adams RP. Composing graphical models with neural networks for structured representations and fast inference. In: Advances in Neural Information Processing Systems. vol. 29; 2016. .
15. Barger Z, Frye CG, Liu D, Dan Y, Bouchard KE. Robust, automated sleep scoring by a compact neural network with distributional shift correction. PLOS ONE. 2019;14(12):e0224642. doi:10.1371/journal.pone.0224642.
16. Grieger N, et al. Automated scoring of pre-REM sleep in mice with deep learning. Scientific Reports. 2021. Supervised pre-REM transitions; verify full citation from publisher before submission.
17. Ebberts J, Heymann J, Drude L, et al. Hidden markov model variational autoencoder for acoustic unit discovery. In: Interspeech; 2017. p. 488-92.